

# Categorical Attributes

---

# Categorical Attributes

Let's switch to a different dataset (a toy one)



- We want to train a model to choose whether to go out and play
- ...Based on weather conditions

# Loading the Data

The dataset is in the `weather.csv` file from the `data` folder

```
In [19]: data = pd.read_csv(os.path.join('.', 'data', 'weather.csv'), sep=',')  
data.head()
```

Out [19]:

	outlook	temperature	humidity	windy	play
0	sunny	85	85	False	no
1	sunny	80	90	True	no
2	overcast	83	86	False	yes
3	rainy	70	96	False	yes
4	rainy	68	80	False	yes

- Several attributes **do not have a numeric value**
- Instead, their value is discrete with no clear ordering, i.e. **categorical**

**We need a **numeric encoding** to handle this data with linear models**

# Encoding Binary Attributes

## Binary attributes can be encoded with the values 0 and 1

This is the case for the columns "windy" and "play"

- First, we tell pandas that the columns have a categorical type

```
In [20]: windy = data['windy'].astype('category')  
play = data['play'].astype('category')  
windy.head()
```

```
Out[20]: 0    False  
1     True  
2    False  
3    False  
4    False  
Name: windy, dtype: category  
Categories (2, bool): [False, True]
```

- Categorical data is still displayed as a string
- ...But internally it is encoded as an integer

# Encoding Binary Attributes

Next, we replace the values with their integer code

We will store the results in a copy of the original table

```
In [21]: data2 = data.copy() # We prepare a copy for the numeric encodings
data2['windy'] = windy.cat.codes
data2['play'] = play.cat.codes
data2.head()
```

Out [21]:

	outlook	temperature	humidity	windy	play
0	sunny	85	85	0	0
1	sunny	80	90	1	0
2	overcast	83	86	0	1
3	rainy	70	96	0	1
4	rainy	68	80	0	1

- Now it is apparent that "windy" and "play" have become numbers

# Encoding Discrete Attributes

**We could use the same approach for discrete attribute in general**

E.g. for the attribute "outlook" in our table

- That would yield a numeric **integer** encoding
- ...Which implies an ordering among the values (e.g. rainy < overcast < sunny)
- When no such ranking exists, this is a bad idea

**In these cases, it is better to adopt a one-hot encoding**

- We introduce a column for each value  $v_k$  of the attribute  $x_j$
- The column contains a 1 iff  $x_j = v_k$ , and 0 otherwise

For example, "sunny | sunny | overcast" becomes:

rainy	overcast	sunny
0	0	1
0	0	1
0	1	0

# Encoding Discrete Attributes

We can obtain a one-hot encoding in pandas via the `get_dummies` method

```
In [22]: data3 = pd.get_dummies(data2, columns=['outlook'])  
data3.head()
```

Out [22]:

	temperature	humidity	windy	play	outlook_overcast	outlook_rainy	outlook_sunny
0	85	85	0	0	False	False	True
1	80	90	1	0	False	False	True
2	83	86	0	1	True	False	False
3	70	96	0	1	False	True	False
4	68	80	0	1	False	True	False

- The method by default processes all columns with categorical or object type
  - Strings in csv files are often parsed as "object" columns
- `get_dummies` can also handle the special case of binary variables
  - ...But I wanted to show you how to obtain an integer encoding, too :-)

# Logistic Regression

---



# Logistic Regression

**Our goal is to predict the value of "play", i.e. a categorical attribute**

We say that we are dealing with a classification problem

- This is second type of ML task
- I.e. another broad definition of an ML problem

**Classification problem can be tackled via Linear Models**

...Via a relatively simple modification

- However, even if it looks like a simple mathematical "hack"
- ...The modification has a strong theoretical basis!

We will discuss this topic a bit in this lecture

**Classification and regression have a distinct statistical foundation**

# Logistic Regression

A linear model for classification can be obtained as follows:

- First, we compute the output as usual:

$$g(x; w) = \sum_{j=1} w_j x_j + w_0$$

- ...But then we feed it to a **logistic function**:

$$\frac{1}{1 + e^{-x}}$$


Overall, we obtain:

$$f(x; w) = \frac{1}{1 + e^{-g(x; w)}}$$

# Logistic Regression

The logistic function is a type of **sigmoid** function

```
In [23]: x = np.linspace(-10, 10, 100)
plt.figure(figsize=(14, 3))
plt.plot(x, 1 / (1 + np.exp(-x)))
plt.tight_layout(); plt.grid(':')
```



- Due to its use, this approach is known as **logistic regression**

# Logistic Regression

## Why using the logistic function?

- We can view the model output as a probability distribution
- Specifically, as the probability of the class being "1"

**With this convention, the target can also be interpreted as a probability**

```
In [24]: data3['play'].head()
```

```
Out[24]: 0    0
         1    0
         2    1
         3    1
         4    1
         Name: play, dtype: int8
```

We view:

- $y_i = 0$  as "the probability of the class being 1 is equal to 0"
- $y_i = 1$  as "the probability of the class being 1 is equal to 1"

# Maximum Likelihood Estimation

**This detail is important because it defines how we perform training**

The process relies on a **change of perspective**

- We pretend that our model is a **data generator**
- ...And compute a formula for the **chance of generating the training set**

This formula is called a **likelihood function**

**From this perspective:**

- Training means to change the model parameters  $w$
- ...So that generating the training set is **as likely as possible**

**This approach is known as Maximum Likelihood Estimation**

- We will see how it can be applied to Logistic Regression
- It's going to be hard: if you get lost, try to understand at least the main idea

# Maximum Likelihood Estimation

**If we assume that  $f(x; w)$  is the source of our data**

...Then, when we have (e.g.)  $f(x; w) = 0.7$ :

- We will generate a 1 with 70% chance
- We will generate a 0 with 30% chance

**Now we can measure the chance that the model makes the right guess:**

- If the label is 1, i.e.  $y_i = 1$ 
  - We will generate that with a  $f(x; w)$  probability
- If the label is 0, i.e.  $y_i = 0$ 
  - We will generate that with a  $1 - f(x; w)$  probability

# Likelihood Function

If we repeat for all examples (assuming statistical independence)...

We get the the probability of correctly generating example in each class.

- For all the examples where the class is 1, we get:

$$\prod_{y_i=1} f(x_i; w)$$

- For all the examples where the class is 0, we get:

$$\prod_{y_i=0} (1 - f(x_i; w))$$

**Intuitively:**

- When we have  $y_i = 1$ , we want  $f(x; w)$  to be high
- When we have  $y_i = 0$ , we want  $f(x; w)$  to be low

# Likelihood Function

With another product we get the chance of generating all the training data

$$L(w) = \prod_{y_i=1} f(x_i; w) \prod_{y_i=0} (1 - f(x_i; w))$$

- This is sort of a probability, but is associated to our model, not to the data itself
- ...And it also depends on the parameters  $w$

This is an example of a likelihood function

**We want to train a model that is a likely source for our data**

This means that we can choose the weights by solving:

$$\operatorname{argmax}_w \log L(w)$$

- I.e. to maximize the likelihood of the data
- This often done via Gradient Descent



# Maximum Likelihood Estimation

**MLE is very important in many Machine Learning approaches**

- It provides a **mathematical foundation** for the training process
- It applies to linear regression, too!
- ...Since the MSE can be interpreted in terms of likelihood

**In practice, scikit-learn does all the heavy lifting for us**

...But understanding the main idea is still very useful

- If you feel confused, that's because likelihood is not an easy concept
- ...But it was worth to at least mention in

# Using Logistic Regression

## Using Logistic Regression in scikit-learn is actually easy

We begin by splitting input/output data as usual:

```
In [25]: cols_in = [c for c in data3.columns if c != 'play']  
  
X = data3[cols_in]  
y = data3['play']  
y.head() # We have a table here, but a vector would also work
```

```
Out[25]: 0    0  
        1    0  
        2    1  
        3    1  
        4    1  
        Name: play, dtype: int8
```

Then the training and test set:

```
In [26]: from sklearn.model_selection import train_test_split  
X_tr, X_ts, y_tr, y_ts = train_test_split(X, y, test_size=0.34, random_state=0)
```

# Using Logistic Regression

Then, we build a `LogisticRegression` model

```
In [27]: from sklearn.linear_model import LogisticRegression  
m = LogisticRegression()
```

...And we call the `fit` method as usual:

```
In [28]: m.fit(X_tr, y_tr);
```

Finally, we can obtain our predictions:

```
In [29]: y_pred_tr = m.predict(X_tr)  
y_pred_ts = m.predict(X_ts)
```

# A Better Look at the Predictions

By default, the prediction is the class with the largest probability

```
In [30]: y_pred_tr
```

```
Out[30]: array([0, 1, 0, 0, 1, 0, 0, 1, 1], dtype=int8)
```

- If we are interested in the raw probability values...
- ...We can call the `predict_proba` method:

```
In [31]: y_prob_tr = m.predict_proba(X_tr)  
y_prob_tr[:5]
```

```
Out[31]: array([[0.72678814, 0.27321186],  
                [0.36752431, 0.63247569],  
                [0.77879214, 0.22120786],  
                [0.77593078, 0.22406922],  
                [0.48441664, 0.51558336]])
```

- Scikit-learn gives us the predicted probability of both classes
- Hence, we get two separate columns

# Evaluation

## We can evaluate the results using metrics

There are four basic metrics for binary classification:

- Number of **True Positives**, i.e.  $TP = \sum_{y_i=1} \tilde{f}(x_i; w)$
- Number of **True Negatives**, i.e.  $TN = \sum_{y_i=0} (1 - \tilde{f}(x_i; w))$
- Number of **False Positives**, i.e.  $FP = \sum_{y_i=0} \tilde{f}(x_i; w)$
- Number of **False Negatives**, i.e.  $FN = \sum_{y_i=1} (1 - \tilde{f}(x_i; w))$

In all cases  $\tilde{f}(x_i; w)$  is the most probable class for the example  $x_i$

# Evaluation

From these we can derive a few more complex metrics

The model (binary) **accuracy** is defined as:

$$ACC = \frac{TP + TN}{m}$$

- I.e. the fraction of examples that is **correctly classified**
- The accuracy ranges over the interval  $[0, 1]$

```
In [32]: from sklearn.metrics import accuracy_score
```

```
print(f'Accuracy on the training set: {accuracy_score(y_tr, y_pred_tr):.3}')
```

```
print(f'Accuracy on the test set: {accuracy_score(y_ts, y_pred_ts):.3}')
```

```
Accuracy on the training set: 0.778
```


```
Accuracy on the test set: 0.8
```

# Evaluation

...Or we can plot all basic metrics via a **confusion matrix**

Here's the one for the training set:

```
In [33]: from sklearn.metrics import ConfusionMatrixDisplay  
ConfusionMatrixDisplay.from_estimator(m, X_tr, y_tr, display_labels=play.cat.categories);
```




# Evaluation

...Or we can plot all basic metrics via a **confusion matrix**

...And the one for the test set

```
In [34]: from sklearn.metrics import ConfusionMatrixDisplay  
ConfusionMatrixDisplay.from_estimator(m, X_ts, y_ts, display_labels=play.cat.categories);
```





# Conclusions and Take-Home Messages

- Handling categorical attributes
  - **Binary** attributes can be handled via 0-1 encoding
  - **Ordinal** attributes via an integer encoding
  - **Categorical** attributes via a one-hot encoding
- Logistic regression is a linear model for **classification** tasks
  - The output can be interpreted as a **probability**
- Training for maximum likelihood
  - The **most common** training method in the ML literature
  - Goal: **maximize the estimated probability** of the training data
- Evaluation of classification models
  - Use **metrics** for a compact evaluation
  - ...And a **confusion matrix** to inspect the details